

STAT40780 Data Programming with C (online)

Lab Sheet 7 (Solutions)

Dr Marie Galligan

Summer 2015

This week's lab requires you to write some C++ functions making use of the Rcpp and inline R packages. There are 3 parts to this lab sheet.

1 A simple function with Rcpp

Write a C++ function (inlined in R code) with Rcpp that accepts two integer scalar arguments passed from R, and returns the first value, divided by the second value. Compile, link, and load this function into R using `cxxfunction()`, then call the function from R.

A possible solution

Run the following R script.

divideInt

```
1 library(Rcpp)
2 library(inline)
3
4 #the function body is defined in this R character string
5 body_divideInt <- '
6   int val1 = as<int>(x1);
7   int val2= as<int>(x2);
8   return wrap( static_cast<double>(val1) / val2 );
9 '
10
11 #compile, link and load the C++ function
12 divideInt <- cxxfunction(
13   signature(x1 = "integer", x2 = "integer"),
14   body = body_divideInt,
15   plugin = "Rcpp"
16 )
17
18 #call the compiled C++ function from R
19 divideInt(as.integer(4), as.integer(3))
20 divideInt(as.integer(8), as.integer(2))
```

2 Practice with IntegerVector

With the help of Rcpp, write a C++ version of the `which.min()` function in R, that accepts as input an integer vector from R. Note: you can assume for now that no missing values will be passed to this function.

If you are unsure what the `which.min()` function does, you can get help on this in R by running

```
?which.min
```

A possible solution

The function below accepts as input an integer vector `x` from R, searches for the minimum value, and returns the position of that element in the vector. e.g. if `x = (9, 8, 2, 3)`, the minimum element is 2, and is the 3rd element in the array, hence the function should return 3. If there are ties (more than one minimum), return the index of the first minimum value encountered.

whichMinCpp

```
1  #A C++ equivalent of the which.min function
2  library(Rcpp)
3  library(inline)
4
5  #create the function body as an R character string
6  body_whichMinCpp <- '
7      IntegerVector xx(x); //convert to IntegerVector object
8      int minval = xx[ 0 ]; //current minimum set to first element
9      int min_index = 1; //current minimum is first element
10
11      for( int i = 1 ; i < xx.size(); i++ )
12      {
13          //check if element indexed by i is
14          //less than the current minimum
15          if( xx[ i ] < minval )
16          {
17              //update minimum value
18              minval = xx[ i ];
19              //update position of minimum element
20              min_index = i+1; //because element indexed by i is the (i+1)th element
21          }
22      } //end of for loop
23
24      return wrap( min_index );
25  '
26
27  #compile, link, load
28  whichMinCpp <- cxxfunction(
29      signature(x = "integer"),
30      body = body_whichMinCpp,
31      plugin = "Rcpp"
32  )
33
34  #create some data to test it on
35  x <- sample(1:20, size = 10, replace = TRUE)
36
37  #call the compiled C++ function
38  whichMinCpp( x )
```

3 Bubblesort function

The C++ code inlined in the R script shown below is the implementation of the bubblesort algorithm from Lesson 4.2, which sorts an `IntegerVector` passed from R. Improve this function as follows:

1. On pass k over the vector, the top $k-1$ elements are already sorted. Modify the bubblesort function so that on pass k , it does not pass over the top $k-1$ elements.
2. A vector with n elements might be sorted in less than $n-1$ passes. Improve the bubblesort algorithm so that it does not continue to pass over a vector that is already sorted.

Benchmark your modified version against this one, to compare performance. Note that it is possible to optimize the bubblesort function further. How?

Bubblesort algorithm

```
1  #body of the C++ bubblesort function
2  #stored in an R character string
3  body_bubblesort2 <- '
4  IntegerVector xx = clone(x); //use of clone()
5  int n = xx.size(); //no. of elements
6  int temp; //temporary storage of swap value
7  for( int k = 1; k <= n - 1; k++){ //for pass k
8    //loop over pairs of elements
9    for( int i = 0; i < n - 1; i ++ ){
10      if( xx[ i ] > xx[ i+1 ] ){
11        temp = xx[ i + 1 ];
12        xx[ i + 1 ] = xx[ i ];
13        xx[ i ] = temp;
14      } //end of if
15    } //end of loop over array pairs
16  } //end of loop over passes
17  return(wrap(xx));
18  '
19
20 #compile, link, load
21 bubblesort2 <- cxxfunction(signature( x = "integer" ),
22                             body = body_bubblesort2,
23                             plugin = "Rcpp")
24
25 #create an R integer vector to input to bubblesort
26 x2 <- as.integer( sample(1:100, size = 100, replace = FALSE) )
27
28 #call
29 bubblesort2(x2) #returns sorted x2
30 x2 #original x2 is not sorted
```

A possible solution

A possible solution is shown here, where modifications 1 and 2 have been made. On Blackboard, I'm uploading some notes on the first modification.

Modified bubblesort algorithm

```
1 #MODIFICATION 1: On pass k, this algorithm does not pass over the top
2 #k-1 elements of the vector
3 #MODIFICATION 2: the algorithm has been modified so that it stops passing
4 #over the vector if no swaps were made during the previous pass
5 body_bubblesortOpt2 <- '
6
7 IntegerVector xx = clone(x); //use of clone()
8 int n = xx.size(); //no. of elements
9 int temp; //temporary storage of swap value
10 int wasSwap = 1; //to test whether a swap was made on pass k
11 //set wasSwap to 0? if no swap was made, 1 otherwise
12
13 //for pass k, continue if k <= n-1
14 //AND if a swap was made on previous pass
15 for( int k = 1; (k <= n - 1) && (wasSwap == 1); k++){
16 wasSwap = 0; //so far, no swap has been made on this pass, so set to zero
17
18 //loop over pairs of elements
19 //modified i < n-1 to i < n-k to
20 //avoid passing over elements that
21 //are already sorted
22 for( int i = 0; i < n - k; i ++ ){
23 if( xx[ i ] > xx[ i+1 ] ){
24 temp = xx[ i + 1 ];
25 xx[ i + 1 ] = xx[ i ];
26 xx[ i ] = temp;
27 wasSwap = 1; // indicate that a swap was made
28 }
29
30 } //end of loop over array pairs
31 } //end of loop over passes
32 return(wrap(xx));
33 ,
34
35 #compile, link, load
36 bubblesortOpt2 <- cxxfunction(signature( x = "integer" ),
37                               body = body_bubblesortOpt2,
38                               plugin = "Rcpp")
39
40 #create some data
41 x <- as.integer( sample(1:100, size = 100, replace = FALSE) )
42 #call the modified bubblesort function
43 sorted_x <- bubblesortOpt2(x) #sorts x
44 sorted_x #sorted_x is sorted
45 x #x is not sorted
```

Benchmarking bubblesort

```
1
2 x <- as.integer( sample(1:100, size = 100, replace = FALSE) )
3 benchmark(bubblesort2(x), bubblesortOpt2(x),
4           order = "relative",
5           replications = 1000)
```

Optimizing bubblesort further: On pass k over the vector, **at least** the top $k-1$ elements will already be sorted - but it is possible that more than $k-1$ top elements will already be sorted. After a pass over the array, all elements after the last swap during that pass are already sorted. So, if after a pass over a vector with 10 elements, the 3th and 4th elements were swapped, but none above this were swapped, then there would be no need to check the top 6 elements on subsequent passes over the array.